

MySQL Replication

Seminario de Desarrollo de Software - Casa Central.

Maximiliano Osorio
mosorio@inf.utfsm.cl

Universidad Técnica Federico Santa María

28 de noviembre de 2017

Replication events

- Statement based – in which case these are write queries
- Row based – in this case these are changes to records, sort of row diffs if you will

On the master

- For replication to work, first of all master needs to be writing replication events to a special log called binary log.
- This is usually very lightweight activity (assuming events are not synchronized to disk), because writes are buffered and because they are sequential.
- The binary log file stores data that replication slave will be reading later.

On the master

- Whenever a replication slave connects to a master, master creates a new thread for the connection
- And then it does whatever the client – replication slave in this case – asks
 1. feeding replication slave with events from the binary log
 2. notifying slave about newly written events to its binary log

On the master

- Slaves that are up to date will mostly be reading events that are still cached in OS cache on the master, so there is not going to be any physical disk reads on the master in order to feed binary log events
- However, when you connect a replication slave that is few hours or even days behind, it will initially start reading binary logs that were written hours or days ago.

On the replica

When you start replication, two threads are started on the slave:

- IO thread: This process called IO thread connects to a master, reads binary log events from the master, they come in and just copies them over to a local log file called relay log.
- SQL thread – reads events from a relay log stored locally on the replication slave (the file that was written by IO thread) and then applies them as fast as possible.

If you want to see where IO thread currently is, check the following in

```
show slave status \G:
```

show slave status

- `Master_Log_File` – last file copied from the master (most of the time it would be the same as last binary log written by a master)
- `Read_Master_Log_Pos` – binary log from master is copied over to the relay log on the slave up until this position.

And then you can compare it to the output of `show master status \G` from the master.

SQL thread

reads events from a relay log stored locally on the replication slave (the file that was written by IO thread) and then applies them as fast as possible.

- `Relay_Master_Log_File` – binary log from master, that SQL thread is "working on" (in reality it is working on relay log, so it's just a convenient way to display information)
- `Exec_Master_Log_Pos` – which position from master binary log is being executed by SQL thread.

Replication lag

First thing you want to know is which of the two replication threads is behind.

- Most of the time it will be the SQL thread
- 1. IO bound
- 2. CPU bound

vmstat

```
[root@smt1 mysql]# vmstat 10
```

procs		-----memory-----				---swap--		-----io-----		--system--		-----cpu-----				
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
13	0	0	880156	40060	2140824	0	0	12	750	895	1045	32	8	54	6	0
0	0	0	877116	40060	2141312	0	0	0	1783	2185	23112	44	10	41	5	0
15	0	0	872648	40068	2141960	0	0	0	1747	2204	25743	41	11	46	2	0
0	0	0	868056	40068	2142604	0	0	0	1803	2164	26224	40	11	44	5	0
17	1	0	863216	40068	2143160	0	0	0	1875	1948	23020	36	9	50	5	0
0	0	0	858384	40168	2143656	0	0	0	1063	1855	21116	32	9	45	14	0
23	0	0	855848	40176	2144232	0	0	0	1755	2036	23181	36	10	48	6	0
49	0	0	851248	40184	2144648	0	0	0	1679	2313	22832	45	10	40	5	0
10	0	0	846292	40192	2145248	0	0	0	1911	1980	23185	36	9	50	4	0
0	0	0	844260	40196	2145868	0	0	0	1757	2152	26387	39	11	45	5	0
0	3	0	839296	40196	2146520	0	0	0	1439	2104	25096	38	10	50	1	0

show processlist

```
mysql> show processlist;
```

-----+-----+-----+-----+-----+-----+-----+-----+-----+-----									
Id	User	db	Command	Time	State	Info	Rows_sent		
Rows_examined		Rows_read							
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----									
1	root	NULL	Sleep	0		NULL		0	
	0	0							
...									
32	root	sbtest	Execute	0	NULL	COMMIT		0	
	0	0							
33	root	sbtest	Execute	0	NULL	COMMIT		0	
	0	0							
34	root	sbtest	Execute	0	Sorting result	SELECT c from			
sbtest where id between 365260 and 365359				order by c		0		0	
	0								
35	root	sbtest	Execute	0	NULL	COMMIT		0	
	0	0							
36	root	sbtest	Execute	0	NULL	COMMIT		0	
	0	0							
37	root	sbtest	Execute	0	NULL	COMMIT		0	
	0	0							
38	root	sbtest	Execute	0	Writing to net	DELETE from sbtest			
where id=496460						0		1	
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----									

fully ACID mode

```
mysql> select @@innodb_flush_log_at_trx_commit;
+-----+
| @@innodb_flush_log_at_trx_commit |
+-----+
|                                1 |
+-----+
```

References



“Mysql replication.” <https://www.percona.com/resources/mysql-ebooks/mysql-replication>.
(Accessed on 10/10/2017).